

DLCV-HW1

- Student ID: B09611007

Problem 1: Self-Supervised Pre-training for Image Classification

1. (5%) Describe the implementation details of your SSL method for pre-training the ResNet50 backbone. (including but not limited to the name of the SSL method & data augmentation techniques you used, learning rate schedule, optimizer, and batch size setting for this pre-training phase)

- Self-Supervised Pre-training using [Bootstrap Your Own Latent \(BYOL\)](#) (<https://github.com/byol-pytorch/trees/master>)
- BYOL is built on two neural networks, referred to as **online networks** and **target networks**, that interact and learn from each other. From an augmented view of an image, the online network is trained to predict the target network representation of the same image under a different augmented view. At the same time, the target network is updated with a slow-moving average of the online network. ([Reference: https://tsang.medium.com/review-byol-bootstrap-your-own-latent-a-new-approach-to-self-supervised-learning-6f770a624441](https://tsang.medium.com/review-byol-bootstrap-your-own-latent-a-new-approach-to-self-supervised-learning-6f770a624441))

- Backbone: ResNet50

- Data Augmentation

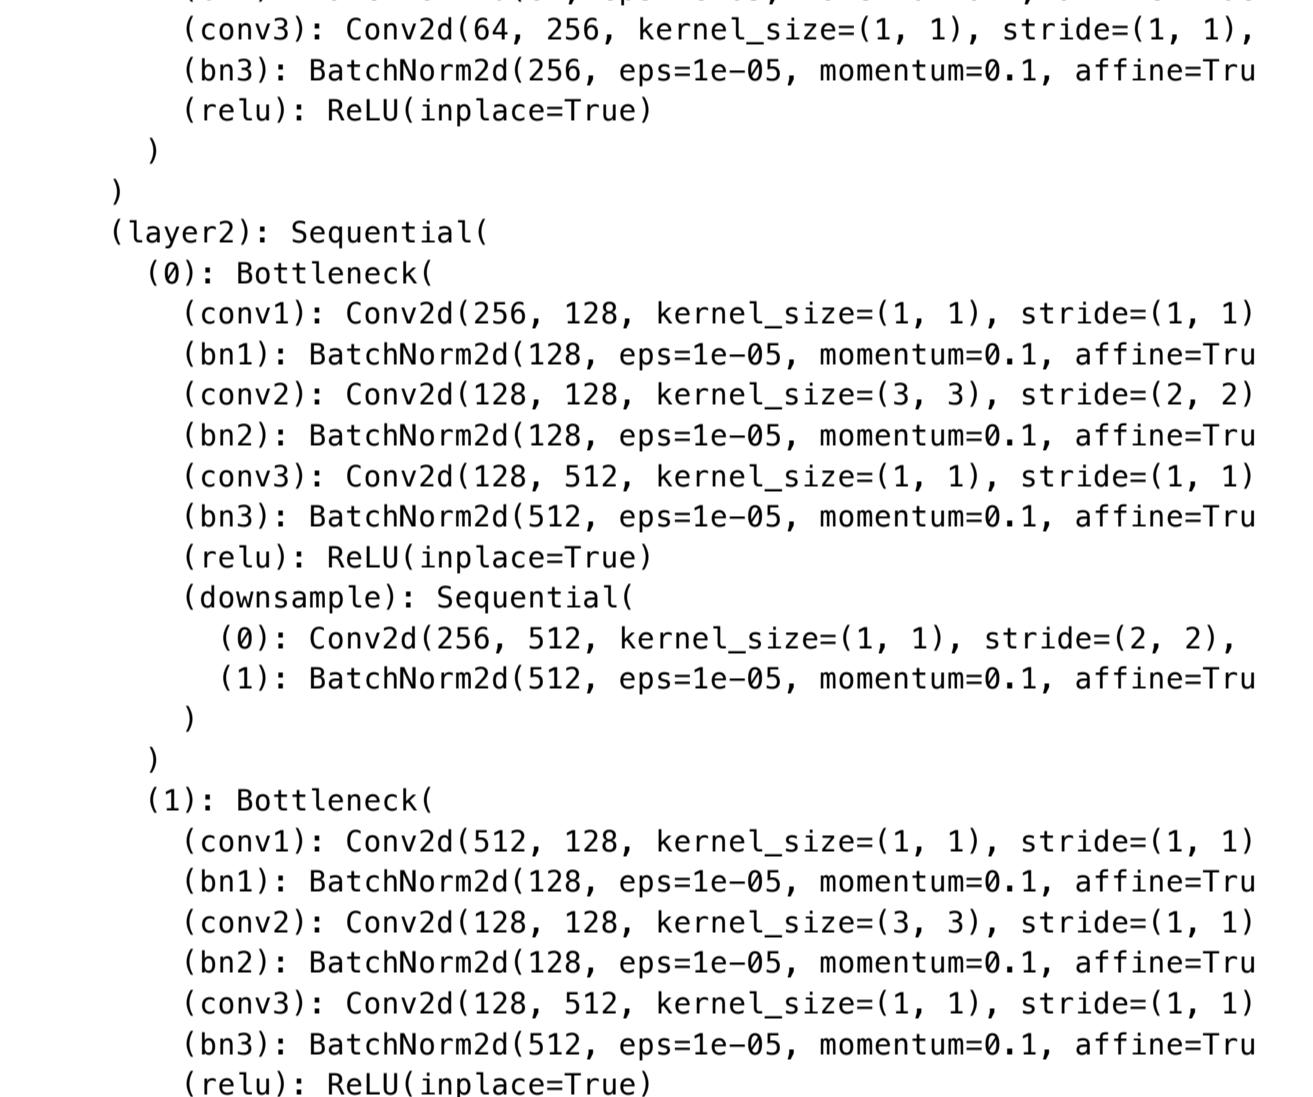
```
RandomApply(
    T.ColorJitter(0.8, 0.8, 0.8, 0.2),
    p=0.3
),  
T.RandomGrayscale(p=0.2),
T.RandomHorizontalFlip(),
RandomApply(
    T.GaussianBlur((3, 3), (1.0, 2.0)),
    p=0.2
),  
T.RandomResizedCrop((128, 128)),
T.Normalize(
    mean=torch.tensor([0.485, 0.456, 0.406]),
    std=torch.tensor([0.229, 0.224, 0.225])),

```

- Pre-training Hyperparams

- epoch: 1000
- batch size: 256
- initial learning rate: 0.00001
- optimizer: Adam
- scheduler: cosine annealing with warmup

- Pre-training Results



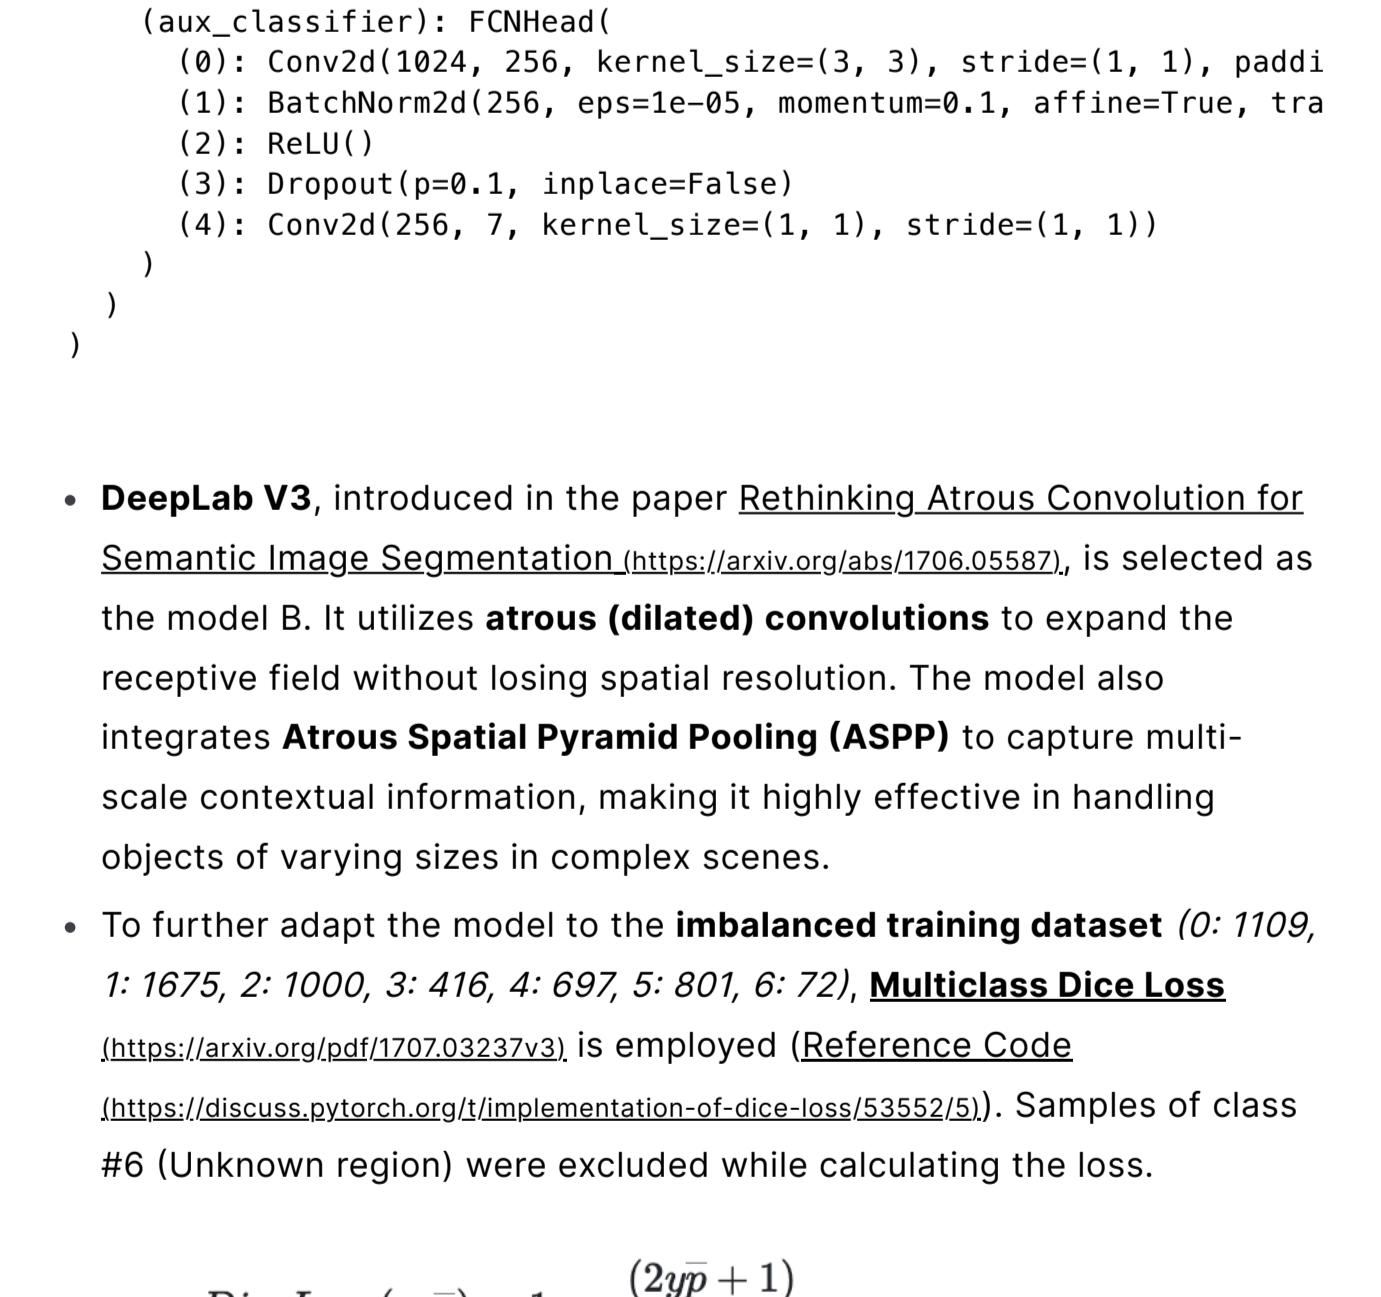
2. (20%) Please conduct the image classification on Office-Home dataset as the downstream task. Also, please complete the following Table, which contains different image classification setting, and discuss/analyze the results.

- **Classifier:** Two FC layer (hidden_size=256) with dropout rate of 0.1

Setting	Pre-training (Mini-ImageNet)	Fine-tuning (Office-Home dataset)	Validation Acc. % (Office-Home dataset)
A	-	Train full model (backbone + classifier)	62.1
B	w/ label	Train full model (backbone + classifier)	61.3
C	w/o label	Train full model (backbone + classifier)	63.5
D	w/ label	Train classifier only	26.0
E	w/o label	Train classifier only	47.9

- **Pre-training with / without labels:** The results show that pretraining without labels (Self-Supervised Learning) outperforms pretraining with labels (Supervised Learning).
- In full model fine-tuning, the model pre-trained without labels achieved higher validation accuracy (C: 63.5% vs. B: 61.3%). Additionally, when only the classifier was fine-tuned, self-supervised pretraining also led to better performance (E: 47.9% vs. D: 26.0%). This suggests that **self-supervised learning helps the model capture more superior performance** even when only fine-tuning the classifier. Conversely, supervised learning may cause the model to **overfit to specific label information**.

3. (5%) Visualize the learned visual representation of setting C on the train set by implementing t-SNE (t-distributed Stochastic Neighbor Embedding) on the output of the second last layer. Depict your visualization from both the first and the last epochs. Briefly explain the results.



- **Epoch 1:** Data points are scattered without clear separation, showing the model hasn't learned meaningful features yet (reflected by low accuracy of 14.5%).
- **Epoch 200:** Data points form clear clusters, suggesting that the model has effectively learned to separate features (leading to higher accuracy of 63.5%).

Problem 2: Semantic Segmentation

1. (3%) Draw the network architecture of your VGG16-FCN32s model (model A).

```
VGG16_FCN32s(  
    features: Sequential(  
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1  
        (1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1  
        (2): Conv2d(64, 128, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1  
        (4): Conv2d(128, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1  
        (6): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (7): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (8): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (9): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (10): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (11): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (12): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (13): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (14): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (15): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (16): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (17): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (18): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (19): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (20): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (21): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (22): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (23): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (24): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (25): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (26): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (27): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (28): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (29): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
        (30): Conv2d(256, 256, kernel_size=2, stride=2, padding=0, dilation=1, ce  
    )  
    (fc6): Sequential(  
        (0): Conv2d(512, 4096, kernel_size=(7, 7), stride=(1, 1))  
        (1): ReLU()  
        (2): Dropout(p=0.5, inplace=False)  
    )  
    (fc7): Sequential(  
        (0): Conv2d(4096, 4096, kernel_size=(1, 1), stride=(1, 1))  
        (1): ReLU()  
        (2): Dropout(p=0.5, inplace=False)  
    )  
    (score_fr): Sequential(  
        (0): Conv2d(4096, 7, kernel_size=(1, 1), stride=(1, 1))  
        (1): ReLU()  
    )  
    (upscore): Sequential(  
        (0): ConvTranspose2d(7, 7, kernel_size=(64, 64), stride=(32, 32))  
    )
```

2. (3%) Draw the network architecture of the improved model (model B) and explain it differs from your VGG16-FCN32s model.

```
DeepLabV3_ResNet50(  
    (model): DeepLabV3()  
        (backbone): IntermediateLayerGetter(  
            (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3  
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, tra  
            (relu): ReLU(inplace=True)  
            (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilati  
            (layer1): Sequential(  
                (0): Bottleneck(  
                    (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (1): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (2): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (3): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (4): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (5): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (6): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (7): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (8): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (9): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (10): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (11): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (12): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (13): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (14): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (15): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),  
                    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv3): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (relu): ReLU(inplace=True)  
                )  
                (16): Bottleneck(  
                    (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1),  
                    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True  
                    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(
```